

August 26, 2026 at 20:31

1. Introduction. This program counts **open** knight’s tours (open Hamiltonian paths) of the $m \times n$ knight graph by a **dual-frontier** transfer that is aggregated with **bucket sorting** rather than a trie, so it parallelizes; and it extracts the **periodic** transfer once the frontier stabilizes, so that far columns are reached by a cheap sparse matrix-vector product (SpMV) instead of re-running the generator.

The method is Knuth’s (as in **DYNHAM/dynahamp**), reimplemented from scratch. Work in the augmented graph $G^+ = G + K_1$: a new **apex** vertex joins every cell, and an open tour of G is a Hamiltonian cycle of G^+ through the apex. Place the cells column by column. After s cells are placed, the **frontier** is the set of not-yet-placed vertices adjacent to a placed one, plus the apex. Each frontier vertex is **bare** (degree 0), **outer** (degree 1, a subpath endpoint), or **inner** (degree 2). The state is that pattern together with the pairing of the outer endpoints. Because it is coded by **relative** positions, not absolute vertex numbers, the code is translation-invariant: the sorted state set repeats with some small period once we are in the bulk, which is exactly what makes the transfer periodic and the SpMV possible.

```
#define _GNU_SOURCE
#define _FILE_OFFSET_BITS 64
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <omp.h>
#include <sys/resource.h>
#include <unistd.h>
#include <pthread.h>

<Types 2>
<Globals 3>
<Subroutines 4>
<Main 32>
```

2. Board and frontier. Cell (r, c) is vertex $v = cm + r$; there are $V = mn$ cells and the apex is V . *frontier_before*(s) lists (sorted) the frontier just before cell s : the apex, s itself (an “extended” frontier, always present), and each future cell $> s$ with an already-placed neighbour.

⟨Types 2⟩ ≡

```
typedef unsigned long long u64;
typedef unsigned __int128 u128;    /* exact build weights (u64 seeds overflow at ml=7) */
typedef unsigned int u32;
```

See also section 10.

This code is used in section 1.

3. #define MAXF 512

#define MAXLEV 40

⟨Globals 3⟩ ≡

```
int m, n, V;
int NB[64 * 64][8], ND[64 * 64];
static const int KR[8] = {-2, -2, -1, -1, 1, 1, 2, 2};
static const int KC[8] = {-1, 1, -2, 2, -2, 2, -1, 1};
int fr[MAXF], ifrb[64 * 64 + 2];
```

See also sections 5, 8, 11, 13, and 30.

This code is used in section 1.

4. All large allocations go through checked wrappers: on failure they print and *exit* gracefully, so a run can never provoke the kernel’s OOM killer.

⟨Subroutines 4⟩ ≡

```

void *xmalloc(size_t n)
{
    void *p = malloc(n);
    if (¬p ∧ n) {
        fprintf(stderr, "OOM: _malloc_%zu_bytes_failed\n", n);
        exit(3);
    }
    return p;
}

void *xrealloc(void *q, size_t n)
{
    void *p = realloc(q, n);
    if (¬p ∧ n) {
        fprintf(stderr, "OOM: _realloc_%zu_bytes_failed\n", n);
        exit(3);
    }
    return p;
}

void *xcalloc(size_t a, size_t b)
{
    void *p = calloc(a, b);
    if (¬p ∧ a ∧ b) {
        fprintf(stderr, "OOM: _calloc_%zu*%zu_failed\n", a, b);
        exit(3);
    }
    return p;
}

```

See also sections 6, 7, 9, 12, 14, 15, 19, 21, 23, 26, 29, and 31.

This code is used in section 1.

5. To bound *physical* memory – the resident set size, which is what actually triggers the kernel OOM killer – a detached watcher thread samples the RSS and *_exit*(3)s gracefully if it exceeds the cap. We deliberately do **not** cap the virtual address space with RLIMIT_AS: glibc arenas, per-thread stacks, and *qsort*’s scratch *mmap* reserve far more virtual than resident memory, so an RLIMIT_AS cap makes those (non-*xmalloc*) allocations fail with a hard crash instead of a clean exit. RSS is the quantity that matters for OOM, and the largest single allocation here (an edge table, a few GB) fits comfortably inside the headroom between the 85%-RAM cap and total RAM, so the 0.5 s sampling never lets resident memory overshoot the machine.

⟨Globals 3⟩ +=

```

long mem_cap_bytes;

```

6. $\langle$ Subroutines 4 $\rangle + \equiv$

```

static long rss_bytes(void)
{
    FILE *f = fopen("/proc/self/statm", "r");
    if (!f) return 0;
    long sz = 0, res = 0;
    if (fscanf(f, "%ld_%ld", &sz, &res) != 2) res = 0;
    fclose(f);
    return res * (long) sysconf(_SC_PAGE_SIZE);
}

static void *mem_watcher(void *arg)
{
    (void) arg;
    for ( ; ; ) {
        long r = rss_bytes();
        if (mem_cap_bytes > 0 & r > mem_cap_bytes) {
            fprintf(stderr,
                "\nMEMORY_CAP: RSS%.1fGB exceeds cap%.1fGB -- exiting"
                "gracefully(raise\n"
                "MEMCAP_GB_to_allow_more)\n", r/1073741824.0, mem_cap_bytes/1073741824.0);
            fflush(stderr);
            _exit(3);
        }
        usleep(300000);
    }
    return 0;
}

/* The binding limit is the tightest of physical RAM and the cgroup memory.max over our
   cgroup-v2 hierarchy: a container/slice cap is often well below physical RAM, and the kernel's
   cgroup OOM killer fires at THAT limit, not the physical one – so a watcher that only knew
   physical RAM would be OOM-killed before it ever tripped. */
static long cgroup_mem_limit(void) { char line [4096] , *p;
    char path[4096] = {0};
    FILE *f = fopen("/proc/self/cgroup", "r");
    if (!f) return -1;
    while ( fgets ( line , sizeof line , f ) ) if ( ( p = strstr ( line , "0: " ) ) )
    {
        p += 3;
        char *nl = strchr(p, '\n');
        if (nl) *nl = 0;
        snprintf(path, sizeof path, "%s", p);
        break;
    }
    fclose(f);
    if (!path[0]) return -1;
    long lim = -1;
    char base[4200];
    snprintf(base, sizeof base, "/sys/fs/cgroup%s", path);
    int up;
    for (up = 0; up < 16; up++) {
        char mm[4300];

```

```

    snprintf(mm, sizeof mm, "%s/memory.max", base);
    FILE *g = fopen(mm, "r");
    if (g) {
        char v[64] = {0};
        if (fgets(v, sizeof v, g) & strncmp(v, "max", 3)) {
            long x = atol(v);
            if (x > 0 & (lim < 0 ∨ x < lim)) lim = x;
        }
        fclose(g);
    }
    if (!strcmp(base, "/sys/fs/cgroup")) break;
    char *s = strrchr(base, '/');
    if (!s) break;
    *s = 0;
    if (strlen(base) < strlen("/sys/fs/cgroup")) break;
}
return lim; } void start_mem_watcher(void)
{
    long ram = sysconf(_SC_PHYS_PAGES) * (long) sysconf(_SC_PAGE_SIZE);
    long cg = cgroup_mem_limit();
    long avail = (cg > 0 & cg < ram) ? cg : ram;    /* tightest binding limit */
    mem_cap_bytes = (long)(avail * 0.85);
    char *e = getenv("MEMCAP_GB");
    if (e) mem_cap_bytes = atol(e) * (1L << 30);
    pthread_t th;
    pthread_create(&th, 0, mem_watcher, 0);
    pthread_detach(th);
    fprintf(stderr, "self_memory_cap: %.0fGBRSS (%s%s; set MEMCAP_GB to override)\n",
        mem_cap_bytes / 1073741824.0, (cg > 0 & cg < ram) ? "cgroup-limited" : "85% of physical",
        (cg > 0 & cg < ram) ? "" : "");
}

```

7. $\langle \text{Subroutines 4} \rangle + \equiv$

```

void build_board(void)
{
  int r, c, k;
  V = m * n;
  for (c = 0; c < n; c++)
    for (r = 0; r < m; r++) {
      int v = c * m + r;
      ND[v] = 0;
      for (k = 0; k < 8; k++) {
        int rr = r + KR[k], cc = c + KC[k];
        if (rr ≥ 0 ∧ rr < m ∧ cc ≥ 0 ∧ cc < n) NB[v][ND[v]++] = cc * m + rr;
      }
    }
}

int frontier_before(int s)
{
  int q = 0, v, u, k;
  static char inF[64 * 64 + 2];
  for (v = 0; v ≤ V; v++) inF[v] = 0;
  inF[V] = 1;
  if (s < V) inF[s] = 1;
  for (v = s + 1; v < V; v++)
    for (k = 0; k < ND[v]; k++) {
      u = NB[v][k];
      if (u < s) {
        inF[v] = 1;
        break;
      }
    }
  for (v = 0; v ≤ V; v++)
    if (inF[v]) {
      fr[q] = v;
      ifrb[v] = q;
      q++;
    }
  return q;
}

```

8. Mate table, derived edges, key, completion. The state is *mate*[] over frontier positions: -2 bare, -1 inner, else the partner position (outer). *add_derived*(*i*, *j*) joins the subpaths ending at *i*, *j*; if they are the two ends of one subpath a cycle forms (flagged in *cycle*, both ends set inner). The key is one byte per position; equal (also column-shifted) patterns give equal keys. A cycle through the apex is a valid open m' -tour iff the covered (inner) cells are the contiguous prefix $\{0..m' - 1\}$ and the rest of the board frontier is bare; *completion_mp* returns that m' .

⟨ Globals 3 ⟩ +≡

int *mate*[MAXF];

int *cycle*;

#pragma ompthreadprivate (*mate*, *cycle*)

9. $\langle \text{Subroutines 4} \rangle + \equiv$

```

int add_derived(int i, int j)
{
    cycle = 0;
    if (mate[i]  $\equiv$  -1  $\vee$  mate[j]  $\equiv$  -1) return 0;
    if (mate[i]  $\equiv$  -2) {
        if (mate[j]  $\equiv$  -2) {
            if (i  $\equiv$  j) {
                cycle = 1;
                return 0;
            }
            mate[i] = j;
            mate[j] = i;
            return 1;
        }
        mate[i] = mate[j];
        mate[mate[j]] = i;
        mate[j] = -1;
        return 1;
    }
    else if (mate[j]  $\equiv$  -2) {
        mate[j] = mate[i];
        mate[mate[i]] = j;
        mate[i] = -1;
        return 1;
    }
    else if (mate[i]  $\neq$  j) {
        mate[mate[i]] = mate[j];
        mate[mate[j]] = mate[i];
        mate[i] = mate[j] = -1;
        return 1;
    }
    mate[i] = mate[j] = -1;
    cycle = 1;
    return 0;
}

int keyof(int q, unsigned char *key)
{
    int i;
    for (i = 0; i < q; i++) key[i] = mate[i]  $\equiv$  -2 ? 0 : mate[i]  $\equiv$  -1 ? 255 : (unsigned char)(1 + mate[i]);
    return q;
}

int completion_mp(int q, int apexpos, int *frn, int s)
{
    int k, mp = s + 1;
    if (mate[apexpos]  $\neq$  -1) return 0;
    k = 0;
    while (k < q  $\wedge$  frn[k]  $\neq$  V  $\wedge$  mate[k]  $\equiv$  -1  $\wedge$  frn[k]  $\equiv$  mp) {
        mp++;
        k++;
    }
}

```

```

for ( ;  $k < q$ ;  $k++$ ) {
  if ( $frn[k] \equiv V$ ) continue;
  if ( $mate[k] \neq -2$ ) return 0;
}
return  $mp$ ;
}

```

10. Bucket aggregation. Each step emits successor records; we sort by key and merge equal keys. *src* carries the emitting state's index, used only while recording the edge tables.

⟨Types 2⟩ +≡

```
typedef struct {
    long off;
    int len;
    u128 w;
    long src;
} Rec;
```

11. ⟨Globals 3⟩ +≡

```
unsigned char *kp;
long kcap, kuse;
Rec *rc;
long nr, rcap;
unsigned char *cb;
#define NTMAX 128
unsigned char *tkp[NTMAX];
long tkuse[NTMAX], tkcap[NTMAX];
Rec *trc[NTMAX];
long tnr[NTMAX], trcap[NTMAX];
int rcmp(const void *A, const void *B)
{
    const Rec *a = A, *b = B;
    int l = a->len < b->len ? a->len : b->len;
    int d = memcmp(cb + a->off, cb + b->off, l);
    return d ? d : a->len - b->len;
}
```

```

12.  ⟨ Subroutines 4 ⟩ +≡
    void emit(unsigned char *key, int len, u128w, long src)
    {
        int t = omp_get_thread_num();
        if (tkuse[t] + len > tkcap[t]) {
            tkcap[t] = tkcap[t] * 2 + len + 65536;
            tkp[t] = xrealloc(tkp[t], tkcap[t]);
        }
        if (tnr[t] ≥ trcap[t]) {
            trcap[t] = trcap[t] * 2 + 65536;
            trc[t] = xrealloc(trc[t], trcap[t] * sizeof(Rec));
        }
        memcpy(tkp[t] + tkuse[t], key, len);
        trc[t][tnr[t]].off = tkuse[t];
        trc[t][tnr[t]].len = len;
        trc[t][tnr[t]].w = w;
        trc[t][tnr[t]].src = src;
        tkuse[t] += len;
        tnr[t]++;
    }
    void merge_pools(void)
    {
        int t;
        long r;
        nr = 0;
        kuse = 0;
        for (t = 0; t < omp_get_max_threads(); t++) {
            for (r = 0; r < tnr[t]; r++) {
                Rec *R = &trc[t][r];
                if (kuse + R->len > kcap) {
                    kcap = kcap * 2 + R->len + 65536;
                    kp = xrealloc(kp, kcap);
                }
                if (nr ≥ rcap) {
                    rcap = rcap * 2 + 65536;
                    rc = xrealloc(rc, rcap * sizeof(Rec));
                }
                memcpy(kp + kuse, tkp[t] + R->off, R->len);
                rc[nr].off = kuse;
                rc[nr].len = R->len;
                rc[nr].w = R->w;
                rc[nr].src = R->src;
                kuse += R->len;
                nr++;
            }
            tnr[t] = 0;
            tkuse[t] = 0;
        }
    }
    int rcmp_r(const void *A, const void *B, void *arg)
    {

```

```

    unsigned char *base = arg;
    const Rec *a = A, *b = B;
    int l = a-len < b-len ? a-len : b-len;
    int d = memcmp(base + a-off, base + b-off, l);
    return d ? d : a-len - b-len;
} /* sort each thread's records in parallel, then merge+reduce across them in one pass (keeping global
    sort order so periodic ids stay consistent), capturing the integer edge table when recording. */
long in_ncur_g; /* Parallel splitter merge: sort each thread's records, sample P-1
    splitter keys that cut the key space into P balanced ranges, binary-search each thread
    for the range boundaries, then merge+reduce each range INDEPENDENTLY in
    parallel, writing to a per-range output buffer. Equal keys never span a range (splitter
    lower_bound is consistent across threads), so the ranges just concatenate. Parallel AND memory - safe :
    no full sorted copy of the records is materialized. */
#define NRNG 256
int keycmp(unsigned char *a, int la, unsigned char *b, int lb)
{
    int l = la < lb ? la : lb;
    int d = memcmp(a, b, l);
    return d ? d : la - lb;
}
typedef struct {
    unsigned char *k;
    int len;
} Samp;
int sampcmp(const void *A, const void *B)
{
    const Samp *a = A, *b = B;
    return keycmp(a-k, a-len, b-k, b-len);
}
static unsigned char *rk[NRNG];
static long rkcap[NRNG], rkuse[NRNG];
static int *rkl[NRNG];
static u128 *rw_[NRNG];
static long rcnt[NRNG], rgcap[NRNG];
static Edge *re[NRNG];
static long rne[NRNG], recap[NRNG];
static long bnd[NTMAX][NRNG+1]; void sort_reduce(int s, int rec) { int NT = omp_get_max_threads(), t, r;
    long i;
#pragma omp parallel for schedule (dynamic, 1)
    for (t = 0; t < NT; t++)
        if (tnr[t]) qsort_r(trc[t], tnr[t], sizeof(Rec), rcmp_r, tkp[t]);
    long tot = 0;
    for (t = 0; t < NT; t++) tot += tnr[t];
    int P = NT * 8;
    if (P > NRNG) P = NRNG;
    if (P < 1) P = 1;
    if ((long) P > tot & tot > 0) P = (int) tot;
    int S = P * 8;
    if ((long) S > tot) S = (int) tot;

```

```

static Samp *samp = 0;
static long sampcap = 0;
if (sampcap < S + 1) {
    sampcap = S + 1;
    samp = xrealloc(samp, sampcap * sizeof(Samp));
}
long sc = 0;
for (t = 0; t < NT ∧ tot > 0; t++) {
    long nt = tnr[t];
    if (nt ≡ 0) continue;
    long take = (long)((double) S * nt / tot);
    if (take < 1) take = 1;
    long j;
    for (j = 0; j < take ∧ sc < S; j++) {
        long idx = (long)((double) j * nt / take);
        if (idx ≥ nt) idx = nt - 1;
        Rec *R = &trc[t][idx];
        samp[sc].k = tkp[t] + R-off;
        samp[sc].len = R-len;
        sc++;
    }
}
S = (int) sc;
qsort(samp, S, sizeof(Samp), sampcmp);
static Samp spl[NRNG];
int np = 0, p;
for (p = 1; p < P; p++) {
    long si = (long)((double) p * S / P);
    if (si ≥ S) si = S - 1;
    if (S > 0) spl[np++] = samp[si];
}
for (t = 0; t < NT; t++) {
    bnd[t][0] = 0;
    bnd[t][P] = tnr[t];
    for (p = 0; p < np; p++) {
        long lo = 0, hi = tnr[t];
        while (lo < hi) {
            long mid = (lo + hi) / 2;
            Rec *R = &trc[t][mid];
            if (keycmp(tkp[t] + R-off, R-len, spl[p].k, spl[p].len) < 0) lo = mid + 1;
            else hi = mid;
        }
        bnd[t][p + 1] = lo;
    }
}
}
#pragma omp parallel for schedule (dynamic, 1)
for (r = 0; r < P; r++) {
    int hp[NTMAX];

```

```

    long hpos[NTMAX];
    int hn = 0, tt;
    for (tt = 0; tt < NT; tt++) hpos[tt] = bnd[tt][r];
#define RLESS(a, b) (
    {
        Rec *Ra = &trc[a][hpos[a]], *Rb = &trc[b][hpos[b]];
        keycmp(tkp[a] + Ra-off, Ra-len, tkp[b] + Rb-off, Rb-len) < 0;
    }
)
    for (tt = 0; tt < NT; tt++)
        if (hpos[tt] < bnd[tt][r + 1]) {
            int c = hn++;
            hp[c] = tt;
            while (c > 0) {
                int pp = (c - 1)/2;
                if (RLESS(hp[c], hp[pp])) {
                    int x = hp[c];
                    hp[c] = hp[pp];
                    hp[pp] = x;
                    c = pp;
                }
                else break;
            }
        }
    }
    rkuse[r] = 0;
    rcnt[r] = 0;
    rne[r] = 0;
    int curlen = -1;
    long curpos = 0;
    u128 sw = 0;
    while (hn > 0) {
        int mt = hp[0];
        Rec *R = &trc[mt][hpos[mt]];
        unsigned char *rkey = tkp[mt] + R-off;
        int rlen = R-len;
        int same = (curlen == rlen & memcmp(rk[r] + curpos, rkey, rlen) == 0);
        if (!same) {
            if (curlen >= 0) {
                if (rcnt[r] >= rgcap[r]) {
                    rgcap[r] = rgcap[r] * 2 + 1024;
                    rkl[r] = xrealloc(rkl[r], rgcap[r] * sizeof(int));
                    rw_[r] = xrealloc(rw_[r], rgcap[r] * sizeof(u128));
                }
                rkl[r][rcnt[r]] = curlen;
                rw_[r][rcnt[r]] = sw;
                rcnt[r]++;
            }
            if (rkuse[r] + rlen > rkcap[r]) {
                rkcap[r] = rkcap[r] * 2 + rlen + 4096;
                rk[r] = xrealloc(rk[r], rkcap[r]);
            }

```

```

    }
    curpos = rkuse[r];
    memcpy(rk[r] + rkuse[r], rkey, rlen);
    rkuse[r] += rlen;
    curlen = rlen;
    sw = 0;
}
sw = red(sw + R-w);
if (rec) {
    if (rne[r] ≥ recap[r]) {
        recap[r] = recap[r] * 2 + 1024;
        re[r] = xrealloc(re[r], recap[r] * sizeof (Edge));
    }
    re[r][rne[r]].src = (int) R-src;
    re[r][rne[r]].dst = (int) rcnt[r];
    re[r][rne[r]].c = 1;
    rne[r]++;
}
hpos[mt]++;
if (hpos[mt] ≥ bnd[mt][r + 1]) hp[0] = hp[--hn];
{
    int c = 0;
    while (1) {
        int l = 2 * c + 1, r2 = 2 * c + 2, sm = c;
        if (l < hn ∧ RLESS(hp[l], hp[sm])) sm = l;
        if (r2 < hn ∧ RLESS(hp[r2], hp[sm])) sm = r2;
        if (sm ≡ c) break;
        int x = hp[c];
        hp[c] = hp[sm];
        hp[sm] = x;
        c = sm;
    }
}
}
if (curlen ≥ 0) {
    if (rcnt[r] ≥ rgcap[r]) {
        rgcap[r] = rgcap[r] * 2 + 1024;
        rkl[r] = xrealloc(rkl[r], rgcap[r] * sizeof (int));
        rw[r] = xrealloc(rw[r], rgcap[r] * sizeof (u128));
    }
    rkl[r][rcnt[r]] = curlen;
    rw[r][rcnt[r]] = sw;
    rcnt[r]++;
}
}

static long base[NRNG + 1], kbase[NRNG + 1], ebase[NRNG + 1];
base[0] = kbase[0] = ebase[0] = 0;
for (r = 0; r < P; r++) {
    base[r + 1] = base[r] + rcnt[r];
    kbase[r + 1] = kbase[r] + rkuse[r];
    ebase[r + 1] = ebase[r] + rne[r];
}

```

```

    }
    ncur = base[P];
    curoff = xrealloc(curoff, (ncur + 1) * sizeof(long));
    curkl = xrealloc(curkl, (ncur + 1) * sizeof(int));
    curw = xrealloc(curw, (ncur + 1) * sizeof (u128));
    curkp = xrealloc(curkp, kbase[P] + 1);
#pragma omp parallel for schedule (dynamic, 1)
    for (r = 0; r < P; r++) {
        long off = kbase[r], j;
        memcpy(curkp + kbase[r], rk[r], rkuse[r]);
        for (j = 0; j < rcnt[r]; j++) {
            curkl[base[r] + j] = rkl[r][j];
            curw[base[r] + j] = rw_[r][j];
            curoff[base[r] + j] = off;
            off += rkl[r][j];
        }
    }
    for (t = 0; t < NT; t++) {
        tnr[t] = 0;
        tkuse[t] = 0;
    }
    if (rec) { int r2;
        static long rne2[NRNG]; /* Coalesce each range's edges IN PLACE and globalize dst – parallel,
            no global consolidation buffer. All edges to a given dst live in one range (its key does), so
            per-range coalescing equals global coalescing; and the ranges hold disjoint, dst-ordered blocks, so
            concatenating them is already globally (dst,src)-sorted – byte-identical to a global sort+coalesce,
            at half the peak edge memory (no eb copy) and cheaper (P small sorts). */
#pragma omp parallel for schedule (dynamic, 1)
        for (r2 = 0; r2 < P; r2++) {
            long pp, o2 = 0;
            if (rne[r2] > 1) qsort(re[r2], rne[r2], sizeof (Edge), ecmp);
            for (pp = 0; pp < rne[r2]; ) {
                long q2 = pp + 1;
                u64 cc = 1;
                while (q2 < rne[r2] ∧ re[r2][q2].src ≡ re[r2][pp].src ∧ re[r2][q2].dst ≡ re[r2][pp].dst) {
                    cc++;
                    q2++;
                }
                re[r2][o2] = re[r2][pp];
                re[r2][o2].c = cc;
                re[r2][o2].dst += base[r2];
                o2++;
                pp = q2;
            }
            rne2[r2] = o2;
        }
        long o = 0;
        for (r2 = 0; r2 < P; r2++) o += rne2[r2];
        nstate[reclev] = in_ncur_g;
        nstate[reclev + 1] = ncur;
        nedge[reclev] = o;
    }

```

```

ncomp[reclv] = reccomp_n;
if (stream_edges) { /* append this level to the open table file (prefix already written) */
    fwrite(&o, sizeof(long), 1, rec_fp);
    for (r2 = 0; r2 < P; r2++)
        if (rne2[r2]) fwrite(re[r2], sizeof(Edge), rne2[r2], rec_fp);
    fwrite(&reccomp_n, sizeof(long), 1, rec_fp);
    fwrite(reccomp_buf, sizeof(Comp), reccomp_n, rec_fp);
    if (ferror(rec_fp)) {
        fprintf(stderr, "disk_write_error_on_%s(out_of_space?)\n", rec_path);
        exit(3);
    }
    edges[reclv] = 0;
    comps[reclv] = 0;
}
else { /* in-process: consolidate into one array for the in-RAM SpMV */
    Edge *eb = xmalloc((o + 1) * sizeof(Edge));
    long enb = 0;
    for (r2 = 0; r2 < P; r2++) {
        long j;
        for (j = 0; j < rne2[r2]; j++) eb[enb++] = re[r2][j];
    }
    edges[reclv] = eb;
    comps[reclv] = xmalloc((reccomp_n + 1) * sizeof(Comp));
    memcpy(comps[reclv], reccomp_buf, reccomp_n * sizeof(Comp));
}
reclv++; } }

```

13. The transfer step. The bucket holds states as (key,weight) over the previous frontier. To place cell s : build the base mate table $bmate$ over the new frontier (survivors carried, newcomers bare) with s in a temporary slot, then bring s to degree 2 by adding its $2 - \deg(s)$ edges to future neighbours or the apex. When rec is set we also record the integer edge table (src index $\rightarrow$ dst index) and the completion contributions for this level.

```

⟨Globals 3⟩ +=
    unsigned char *curkp;
    long *curoff;
    int *curkl;
    u128 *curw;
    long ncur;
    u128 cnt[1 << 16];
    u64 MODP = 0; /* 0 = exact u64; set to a prime for one CRT residue */
    static inline u128 red(u128 x)
    {
        return MODP ? x % MODP : x;
    }
    int qnew, posS, apexnew, STEMP;
    int bmate[MAXF], o2n[MAXF];
#pragma omp threadprivate (bmate, STEMP)
    int recording, reclev;
    typedef struct {
        int src, dst;
        u64 c;
    } Edge;
    typedef struct {
        int src, delta;
        u64 mult;
    } Comp;
    Edge *edges[MAXLEV];
    long nedge[MAXLEV];
    Comp *comps[MAXLEV];
    long ncomp[MAXLEV];
    long nstate[MAXLEV + 1];
    int Plevs;
    int period_g, c0_g, recend_col_g; /* saved for dumping the extracted tables */
    int direct_valid_col_g; /* direct cnt is exact for columns ≤ this */
    int stop_after_record = 0;
    int direct_cov_g = 0;
    int stream_edges = 0;
    FILE *rec_fp = 0;
    char rec_path[4096];
    long nstate_off = 0; /* out-of-core recording */
    int stream_spmv = 0;
    FILE *tbl_fp = 0;
    long lev_edge_off[MAXLEV]; /* out-of-core SpMV (edges streamed from disk) */
    static u64 colfp_g[4096];
    const char *ckpt_path = 0;
    int ckpt_every = 4; /* checkpoint every K columns */
    u128 *seedv;

```

```

long seedn;
Comp *reccomp_buf;
long reccomp_n, reccomp_cap;
int ecmp(const void *A, const void *B)
{
    const Edge *a = A, *b = B;
    if (a→dst ≠ b→dst) return a→dst - b→dst;
    return a→src - b→src;
}

```

14. ⟨Subroutines 4⟩ +≡

```

void build_bmate(int *omate, int qold)
{
    int i;
    for (i = 0; i ≤ qnew; i++) bmate[i] = -2;
    STEMP = qnew;
    for (i = 0; i < qold; i++) {
        int dst = (i ≡ posS) ? STEMP : o2n[i];
        if (dst < 0) continue;
        if (omate[i] ≡ -1) bmate[dst] = -1;
        else if (omate[i] ≥ 0) {
            int op = omate[i];
            int pdst = (op ≡ posS) ? STEMP : o2n[op];
            bmate[dst] = pdst ≥ 0 ? pdst : -2;
        }
    }
}

```

15. $\langle \text{Subroutines 4} \rangle + \equiv$
void *run_step*(**int** *s*, **int** *rec*)
{
 int *i*;
 long *si*;
 long *in_ncur* = *ncur*;
 int *qold* = *frontier_before*(*s*);
 posS = *ifrb*[*s*];
 static int *frolld*[MAXF];
 for (*i* = 0; *i* < *qold*; *i*++) *frolld*[*i*] = *fr*[*i*];
 qnew = *frontier_before*(*s* + 1);
 static int *ifrnew*[64 * 64 + 2], *frnew*[MAXF];
 for (*i* = 0; *i* < *qnew*; *i*++) *frnew*[*i*] = *fr*[*i*];
 for (*i* = 0; *i* ≤ *V*; *i*++) *ifrnew*[*i*] = -1;
 for (*i* = 0; *i* < *qnew*; *i*++) *ifrnew*[*frnew*[*i*]] = *i*;
 apexnew = *ifrnew*[*V*];
 int *nbr*[16], *rr* = 0;
 nbr[*rr*++] = *apexnew*;
 for (*i* = 0; *i* < *ND*[*s*]; *i*++) {
 int *w* = *NB*[*s*][*i*];
 if (*w* > *s* ∧ *ifrnew*[*w*] ≥ 0) *nbr*[*rr*++] = *ifrnew*[*w*];
 }
 for (*i* = 0; *i* < *qold*; *i*++) *o2n*[*i*] = (*frolld*[*i*] ≡ *s*) ? -1 : *ifrnew*[*frolld*[*i*]];
 in_ncur_g = *ncur*;
 recomp_n = 0;
 $\langle \text{Expand every state at cell } s \text{ 16} \rangle$;
 $\langle \text{Sort, reduce, and (if recording) capture the edge table 18} \rangle$;
}

16. Each state expands independently, so the loop runs in parallel (except while recording, when it stays serial to keep the completion log ordered). The transition scratch (*mate*, *bmate*, *cycle*, **STEMP**) is *threadprivate*; each thread emits into its own pool; *cnt* is updated in a critical section.

```

⟨Expand every state at cell s 16⟩ ≡
#pragma omp parallel for schedule (dynamic, 32)
  for (si = 0; si < ncur; si++) {
    int i, a, b, nl, deg, need, mp;
    unsigned char *ok = curkp + curoff[si];
    int okl = curkl[si];
    u128w = curw[si];
    int omate[MAxF];
    unsigned char nk[MAxF];
    for (i = 0; i < okl; i++) {
      int cc = ok[i];
      omate[i] = cc ≡ 0 ? -2 : cc ≡ 255 ? -1 : cc - 1;
    }
    deg = omate[posS] ≡ -2 ? 0 : omate[posS] ≡ -1 ? 2 : 1;
    need = 2 - deg;
    if (need ≡ 0) {
      build_bmate(omate, gold);
      for (i = 0; i < qnew; i++) mate[i] = bmate[i];
      nl = keyof(qnew, nk);
      emit(nk, nl, w, si);
    }
    else if (need ≡ 1) {
      for (a = 0; a < rr; a++) {
        build_bmate(omate, gold);
        for (i = 0; i ≤ qnew; i++) mate[i] = bmate[i];
        if (add_derived(STEMP, nbr[a])) {
          nl = keyof(qnew, nk);
          emit(nk, nl, w, si);
        }
        else if (cycle) {
          mp = completion_mp(qnew, apexnew, frnew, s);
          if (mp) ⟨Credit completion 17⟩;
        }
      }
    }
    else {
      for (a = 0; a < rr; a++)
        for (b = a + 1; b < rr; b++) {
          build_bmate(omate, gold);
          for (i = 0; i ≤ qnew; i++) mate[i] = bmate[i];
          if (¬add_derived(STEMP, nbr[a])) {
            if (cycle) {
              mp = completion_mp(qnew, apexnew, frnew, s);
              if (mp) ⟨Credit completion 17⟩;
            }
            continue;
          }
          if (add_derived(STEMP, nbr[b])) {

```

```

        nl = keyof(qnew, nk);
        emit(nk, nl, w, si);
    }
    else if (cycle) {
        mp = completion_mp(qnew, apexnew, frnew, s);
        if (mp) < Credit completion 17 >;
    }
}
}
}

```

This code is used in section 15.

17. < Credit completion 17 > $\equiv$

```

{
#pragma omp critical
{
    cnt[mp] = red(cnt[mp] + w);
    if (rec) {
        if (recomp_n ≥ recomp_cap) {
            recomp_cap = recomp_cap * 2 + (1 ≪ 20);
            recomp_buf = xrealloc(recomp_buf, recomp_cap * sizeof(Comp));
        }
        recomp_buf[recomp_n].src = si;
        recomp_buf[recomp_n].delta = mp - (s + 1);
        recomp_buf[recomp_n].mult = 1;
        recomp_n++;
    }
}
}

```

This code is used in section 16.

18. The sort and reduce run across the per-thread pools in parallel (see *sort_reduce*); *in_ncur_g* records the input level size for the edge table.

< Sort, reduce, and (if recording) capture the edge table 18 > $\equiv$
sort_reduce(s, rec);

This code is used in section 15.

19. Checkpointing the build. The build (the sweep that discovers the periodic transfer) is the long part; we checkpoint it at every column boundary before recording begins, dumping the current bucket, the running counts, and the fingerprints. A build that dies resumes from the last boundary; the recording columns themselves are few, so if a crash lands inside them we simply restart from the last pre-recording checkpoint.

⟨Subroutines 4⟩ +≡

```

void save_ckpt(int nexts)
{
    if ( $\neg$ ckpt_path) return;
    FILE *f = fopen(ckpt_path, "wb");
    fwrite(&m, sizeof(int), 1, f);
    fwrite(&nexts, sizeof(int), 1, f);
    fwrite(cnt, sizeof (u128), 1  $\ll$  16, f);
    fwrite(colfp_g, sizeof(u64), 4096, f);
    fwrite(&ncur, sizeof(long), 1, f);
    fwrite(&kuse, sizeof(long), 1, f);
    fwrite(curoff, sizeof(long), ncur, f);
    fwrite(curkl, sizeof(int), ncur, f);
    fwrite(curw, sizeof (u128), ncur, f);
    {
        long kb = 0, i;
        for (i = 0; i < ncur; i++) kb += curkl[i];
        fwrite(&kb, sizeof(long), 1, f);
        fwrite(curkp, 1, kb, f);
    }
    fclose(f);
}

int load_ckpt(const char *path)
{
    FILE *f = fopen(path, "rb");
    if ( $\neg$ f) return -1;
    int nexts, mm;
    if (fread(&mm, sizeof(int), 1, f)  $\neq$  1) return -1;
    m = mm;
    fread(&nexts, sizeof(int), 1, f);
    fread(cnt, sizeof (u128), 1  $\ll$  16, f);
    fread(colfp_g, sizeof(u64), 4096, f);
    fread(&ncur, sizeof(long), 1, f);
    fread(&kuse, sizeof(long), 1, f);
    curoff = xrealloc(curoff, (ncur + 1) * sizeof(long));
    curkl = xrealloc(curkl, (ncur + 1) * sizeof(int));
    curw = xrealloc(curw, (ncur + 1) * sizeof(u64));
    fread(curoff, sizeof(long), ncur, f);
    fread(curkl, sizeof(int), ncur, f);
    fread(curw, sizeof (u128), ncur, f);
    {
        long kb;
        fread(&kb, sizeof(long), 1, f);
        curkp = xrealloc(curkp, kb + 1);
        fread(curkp, 1, kb, f);
    }
}

```

```
    }  
    fclose(f);  
    fprintf(stderr, "resumed build from %s at cell %d (col %d)\n", path, nexts, nexts/mm);  
    return nexts;  
}
```

20.

21. Dumping and reloading the extracted tables. The expensive part is the build that discovers and records the periodic transfer. Once recorded, the tables (edge tables, completions, seed vector, level sizes, period, and the directly-counted columns up to *recend_col_g*) are small and self-contained, so we dump them to disk. A later run reloads them and does only the cheap SpMV — and a build that dies can be re-run without touching the SpMV.

⟨Subroutines 4⟩ +≡

```

void dump_tables(const char *path)
{
    int L;
    if (stream_edges ∧ rec_fp) { /* The prefix (hdr, seed, nstate placeholder) and every level
        are already on disk (written during recording). Append cnt, then patch the two fields
        not known until now: hdr[5]=direct_valid_col_g and the real nstate[].Nocopy, no second file —
        — so the table needs disk for its own size only. */
        int nc = (direct_valid_col_g + 1) * m;
        fwrite(&nc, sizeof(int), 1, rec_fp);
        fwrite(cnt, sizeof (u128), nc, rec_fp);
        fflush(rec_fp);
        fseek(rec_fp, 5L * sizeof(int), SEEK_SET);
        fwrite(&direct_valid_col_g, sizeof(int), 1, rec_fp);
        fseek(rec_fp, nstate_off, SEEK_SET);
        fwrite(nstate, sizeof(long), Plevs + 1, rec_fp);
        if (ferror(rec_fp)) {
            fprintf(stderr, "disk_write_error_on_%s(out_of_space?)\n", path);
            exit(3);
        }
        fclose(rec_fp);
        rec_fp = 0;
        fprintf(stderr, "dumped_tables_to_%s(period_%d,c0_%d,%d_levels)\n", path, period_g, c0_g,
            Plevs);
        return;
    }

    FILE *f = fopen(path, "wb"); /* non-streaming fallback (edges held in RAM) */
    int hdr[6] = {m, period_g, c0_g, Plevs, recend_col_g, direct_valid_col_g};
    fwrite(hdr, sizeof(int), 6, f);
    fwrite(&seedn, sizeof(long), 1, f);
    fwrite(seedv, sizeof (u128), seedn, f);
    fwrite(nstate, sizeof(long), Plevs + 1, f);
    for (L = 0; L < Plevs; L++) {
        fwrite(&nedge[L], sizeof(long), 1, f);
        fwrite(edges[L], sizeof(Edge), nedge[L], f);
        fwrite(&ncomp[L], sizeof(long), 1, f);
        fwrite(comps[L], sizeof(Comp), ncomp[L], f);
    }
    int nc = (direct_valid_col_g + 1) * m;
    fwrite(&nc, sizeof(int), 1, f);
    fwrite(cnt, sizeof (u128), nc, f);
    fclose(f);
    fprintf(stderr, "dumped_tables_to_%s(period_%d,c0_%d,%d_levels)\n", path, period_g, c0_g,
        Plevs);
}

```

```

int load_tables(const char *path)
{
    FILE *f = fopen(path, "rb");
    if (!f) return 0;
    int L, hdr[6];
    if (fread(hdr, sizeof(int), 6, f) != 6) return 0;
    m = hdr[0];
    period_g = hdr[1];
    c0_g = hdr[2];
    Plevs = hdr[3];
    recend_col_g = hdr[4];
    direct_valid_col_g = hdr[5];
    fread(&seedn, sizeof(long), 1, f);
    seedv = xmalloc(seedn * sizeof(u128));
    fread(seedv, sizeof(u128), seedn, f);
    fread(nstate, sizeof(long), Plevs + 1, f);
    /* Stream edges from disk (tables can be in RAM, e.g. 100GB for m=7): keep the small comps[] in
       RAM, remember each level's edge offset, skip the edges. */
    for (L = 0; L < Plevs; L++) {
        fread(&nedge[L], sizeof(long), 1, f);
        lev_edge_off[L] = ftello(f);
        fseeko(f, (off_t)nedge[L] * sizeof(Edge), SEEK_CUR);
        fread(&ncomp[L], sizeof(long), 1, f);
        comps[L] = xmalloc((ncomp[L] + 1) * sizeof(Comp));
        fread(comps[L], sizeof(Comp), ncomp[L], f);
        edges[L] = 0;
    }
    int nc;
    fread(&nc, sizeof(int), 1, f);
    memset(cnt, 0, sizeof(cnt));
    fread(cnt, sizeof(u128), nc, f);
    tbl_fp = f;
    stream_spmv = 1;
    return 1;
} /* keep file open for edge streaming */

```

22.

23. Driver: build, extract the period, then SpMV. Sweep an $m \times W_b$ strip; when the boundary key-set (fingerprinted) repeats with period $p \in \{1, 2\}$ at column c_0 , record the next p columns' edge tables and capture the seed vector. Then iterate the SpMV to reach column N : at each recorded level apply the edge table to the state vector and add the level's completions to *cnt2*. The two agree on every column the SpMV fully covers ($c \geq c_0 + 3$), and the SpMV alone yields the columns past W_b .

⟨Subroutines 4⟩ +≡

```

int resume_from = 0;
void seed_bucket(void)
{
    int i;
    int q = frontier_before(0);
    for (i = 0; i < q; i++) mate[i] = -2;
    unsigned char key[MAXF];
    int kl = keyof(q, key);
    curkp = xmalloc(1 << 20);
    curoff = xmalloc(sizeof(long));
    curkl = xmalloc(sizeof(int));
    curw = xmalloc(sizeof(u64));
    memcpy(curkp, key, kl);
    curoff[0] = 0;
    curkl[0] = kl;
    curw[0] = 1;
    ncur = 1;
}
u64 cnt2[1 << 16];
int run_periodic(int mm, int Wb, int Next)
{
    int i, s, c;
    m = mm;
    n = Wb;
    build_board();
    memset(cnt2, 0, sizeof(cnt2));
    recording = 0;
    reclev = 0;
    c0_g = -1;
    recend_col_g = -1;
    direct_valid_col_g = -1;
    Plevs = 0;
    seedn = 0;
    if (resume_from <= 0) {
        memset(cnt, 0, sizeof(cnt));
        seed_bucket();
    } /* fresh; else state is loaded */
    int c0 = -1, period = 0, recstart = -1, recend = -1, seam_break = -1;
    int cand_p = 0, cand_c = -1; /* pending (unconfirmed) period candidate */
    for (s = (resume_from > 0 ? resume_from : 0); s < V; s++) {
        if (recording & s == recstart) {
            seedn = ncur;
            seedv = xmalloc(ncur * sizeof(u128));
            for (i = 0; i < ncur; i++) seedv[i] = curw[i];

```

```

    nstate[0] = ncur;
    reclev = 0;
    if (stream_edges) {
        /* write the table's fixed prefix now; levels stream in after; patch at dump */
        rec_fp = fopen(rec_path, "wb");
        if (!rec_fp) {
            fprintf(stderr, "cannot open %s\n", rec_path);
            exit(3);
        }
        int hdr[6] = {m, period_g, c0_g, Plevs, recend_col_g, 0}; /* direct_valid_col_g patched at dump */
        fwrite(hdr, sizeof(int), 6, rec_fp);
        fwrite(&seedn, sizeof(long), 1, rec_fp);
        fwrite(seedv, sizeof(u128), seedn, rec_fp);
        nstate_off = ftell(rec_fp);
        {
            long z[MAXLEV + 1];
            int L;
            for (L = 0; L ≤ Plevs; L++) z[L] = 0;
            fwrite(z, sizeof(long), Plevs + 1, rec_fp);
        }
        if (ferror(rec_fp)) {
            fprintf(stderr, "disk write error on %s (out of space?)\n", rec_path);
            exit(3);
        }
    }
}
run_step(s, recording ∧ s ≥ recstart ∧ s ≤ recend);
if (s % m ≡ m - 1) {
    c = s/m;
    u64 h = 1469598103934665603ULL;
    long t;
    for (t = 0; t < ncur; t++) {
        unsigned char *kk = curkp + curoff[t];
        int L = curkl[t], z;
        for (z = 0; z < L; z++) {
            h ⊕= kk[z];
            h *= 1099511628211ULL;
        }
        h ⊕= #9e;
        h *= 1099511628211ULL;
    }
    colfp_g[c] = h;
    ⟨Detect and confirm the periodic column 24⟩
}
if (s % m ≡ m - 1 ∧ getenv("DBG"))
    fprintf(stderr, "col %d ncur=%ld rec=%d\n", s/m, ncur, recording);
if (!recording ∧ s % m ≡ m - 1 ∧ (s/m) % ckpt_every ≡ 0) save_ckpt(s + 1);
if (recording ∧ s ≡ seam_break ∧ stop_after_record) break;
}
⟨Finalize direct coverage and verify periodicity 25⟩
direct_cov_g = stop_after_record ? direct_valid_col_g : n;

```

```

    if ( $\neg$ stream_edges) spmv_run(Next, 1);    /* streamed edges are freed; SpMV via a separate 'run' */
    return c0;
}

```

24. We commit to a period only once the boundary fingerprint has *repeated*: a lone match $colfp[c] \equiv colfp[c - p]$ can be a coincidence of the near-boundary columns, so we hold it as a candidate and require it to persist one full period further before recording. That guarantees the recorded columns sit in the bulk, far enough from the strip's right edge that the transfer is stationary; a strip too narrow to contain a confirmed period never records, which the finalizer below turns into a clear diagnostic instead of a corrupt table.

(Detect and confirm the periodic column 24) $\equiv$

```

if (c0 < 0) {
    if (cand_p  $\equiv$  0) {
        if ( $c \geq 1 \wedge colfp\_g[c] \equiv colfp\_g[c - 1]$ ) {
            cand_p = 1;
            cand_c = c;
        }
        else if ( $c \geq 2 \wedge colfp\_g[c] \equiv colfp\_g[c - 2]$ ) {
            cand_p = 2;
            cand_c = c;
        }
    }
    else if ( $colfp\_g[c] \equiv colfp\_g[c - cand\_p]$ ) {
        if ( $c - cand\_c \geq cand\_p \wedge c + cand\_p + 3 \leq n$ ) {
            /* confirmed, with bulk room for recording+seam */
            period = cand_p;
            c0 = c;
            recstart = (c0 + 1) * m;
            recend = (c0 + 1 + period) * m - 1;
            Plevs = period * m;
            recording = 1;
            period_g = period;
            c0_g = c0;
            recend_col_g = c0 + period;
            seam_break = recend + 3 * m;    /* sweep a few extra columns (direct seam) */
            fprintf(stderr, "stable_col_%d, %d (confirmed)\n", c0, period);
        }
    }
    else {
        cand_p = 0;    /* candidate broke: restart search here */
        if ( $c \geq 1 \wedge colfp\_g[c] \equiv colfp\_g[c - 1]$ ) {
            cand_p = 1;
            cand_c = c;
        }
        else if ( $c \geq 2 \wedge colfp\_g[c] \equiv colfp\_g[c - 2]$ ) {
            cand_p = 2;
            cand_c = c;
        }
    }
}
}

```

This code is used in section 23.

25. The direct sweep counts every column it reaches, exact once that column's completions have all closed (a lag of a couple knight-moves). The SpMV cannot reconstruct completions that close inside the recorded/seed columns, so its *first* extrapolated column, *recend* + 1, comes out short. We therefore trust the direct count one column past the recorded region — the *seam_break* sweep ran far enough for its completions to close — and hand over to the SpMV only from *recend* + 2, exactly where the cross-check begins. If a period was recorded we also assert the transfer is stationary; a mismatch means the strip was too narrow and the recording is boundary-contaminated.

⟨Finalize direct coverage and verify periodicity 25⟩ ≡

```
{
  int swept = (s ≥ V) ? n : s/m;
  if (c0 < 0) {
    direct_valid_col_g = swept > 0 ? swept - 1 : 0;    /* no period: direct only */
    fprintf(stderr, "no_confirmed_period_within_strip_W_b=%d: direct_counts_only\
      \n"(increase_W_b_to_record_the_periodic_transfer)\n", n);
  }
  else {
    direct_valid_col_g = (swept ≥ recend_col_g + 3) ? recend_col_g + 1 : recend_col_g;
    if (nstate[0] ≠ nstate[Plevs]) {
      fprintf(stderr,
        "ERROR: recorded_transfer_not_stationary_(nstate[0]=%ld \"nstate[Plevs]=%ld);\
        \nstrip_W_b=%d_too_narrow--recording_columns_are_\"boundary-contaminat\
        ed.Rebuild_with_a_larger_W_b.\n", nstate[0], nstate[Plevs], n);
      exit(4);
    }
  }
}
```

This code is used in section 23.

26. $\langle$ Subroutines 4 $\rangle + \equiv$ /* SpMV from the (built or loaded) tables out to column *Nto*. *crosscheck* compares against the direct *cnt* where both are valid. Uses *cnt2* scratch. */

```

int build_extract(int mm, int Wb)
{
    return run_periodic(mm, Wb, Wb);
}

void spmv_run(int Nto, int crosscheck)
{
    int i;
    memset(cnt2, 0, sizeof (cnt2));
     $\langle$  Iterate the SpMV out to column N 27  $\rangle$ ;
    {
        int good = 1, cc;
        if (crosscheck)
            for (cc = c0_g + 3; cc  $\leq$  Nto  $\wedge$  cc  $\leq$  direct_cov_g; cc++)
                if ((u64) cnt[cc * m]  $\neq$  cnt2[cc * m]) {
                    good = 0;
                    fprintf(stderr, "MISMATCH_c=%d_d_direct=%llu_spmv=%llu\n", cc, (unsigned long
                        long)(u64) cnt[cc * m], (unsigned long long) cnt2[cc * m]);
                }
        for (cc = 1; cc  $\leq$  Nto; cc++) {
            u64 val = cc  $\leq$  direct_valid_col_g ? (u64) red(cnt[cc * m]) : cnt2[cc * m];
            if (val) printf("open_%dx%d=%llu%s\n", m, cc, val, cc > direct_valid_col_g ? "_(SpMV)" : "");
        }
        if (crosscheck) fprintf(stderr, "%s\n", good ? "SpMV_matches_direct" : "SpMV_MISMATCH");
    }
}

```

```

27.  $\langle \text{Iterate the SpMV out to column } N \text{ 27} \rangle \equiv$ 
  if ( $c0\_g \geq 0 \wedge Plevs > 0$ ) { u64 *v = xmalloc(nstate[0] * sizeof(u64));
  {
    long i2;
    for (i2 = 0; i2 < nstate[0]; i2++) v[i2] = (u64)(MODP ? seedv[i2] % MODP : seedv[i2]);
  }
  int basecol = c0_g + 1; while (basecol ≤ Nto) { int L; for (L = 0; L < Plevs; L++) { int
    abscol = basecol + L/m, substep = L % m;
  long e;
  for (e = 0; e < ncomp[L]; e++) {
    int idx = abscol * m + substep + 1 + comps[L][e].delta;
    cnt2[idx] = red(cnt2[idx] + v[comps[L][e].src] * comps[L][e].mult);
  }
  u64 *vn = xcalloc(nstate[L + 1], sizeof(u64)); { long ne = nedge[L];
  if (stream_spmv) { /* out-of-core: stream this level's edges from disk in chunks */
    fseeko(tbl_fp, lev_edge_off[L], SEEK_SET);
    static Edge *buf = 0;
    static long bufcap = 0;
    long CH = 1_L ≪ 20;
    if (bufcap < CH) {
      bufcap = CH;
      buf = xrealloc(buf, bufcap * sizeof(Edge));
    }
    long done = 0;
    while (done < ne) {
      long chunk = ne - done;
      if (chunk > CH) chunk = CH;
      if ((long)fread(buf, sizeof(Edge), chunk, tbl_fp) ≠ chunk) {
        fprintf(stderr, "table_edge_read_error\n");
        exit(3);
      }
      long e2;
      for (e2 = 0; e2 < chunk; e2++) {
        int d = buf[e2].dst;
        vn[d] = red(vn[d] + v[buf[e2].src] * buf[e2].c);
      }
      done += chunk;
    }
  }
  }
  else { /* in-RAM: dst-partitioned parallel apply */
  int NT = omp_get_max_threads();
  static long spl[NTMAX + 1];
  spl[0] = 0;
  spl[NT] = ne;
  for (tt = 1; tt < NT; tt++) {
    long g = (long)((double)tt * ne / NT);
    while (g > 0 ∧ g < ne ∧ edges[L][g].dst ≡ edges[L][g - 1].dst) g++;
    spl[tt] = g;
  }
  }

```

```

    }
    #pragma omp parallel for schedule (dynamic,1)
    for (tt = 0; tt < NT; tt++) {
        long e2;
        for (e2 = spl[tt]; e2 < spl[tt + 1]; e2++) {
            int d = edges[L][e2].dst;
            vn[d] = red(vn[d] + v[edges[L][e2].src] * edges[L][e2].c);
        }
    }
    } } free(v);
    v = vn; } basecol += period_g; } free(v); }

```

This code is used in section 26.

28. $\langle$ Report and cross-check 28 $\rangle \equiv$

```

{
    int good = 1, cc;
    for (cc = c0 + 3; cc ≤ Next ∧ cc ≤ Wb; cc++)
        if (cnt[cc * m] ≠ cnt2[cc * m]) {
            good = 0;
            fprintf(stderr, "MISMATCH_c=%d_direct=%llu_spmv=%llu\n", cc, cnt[cc * m], cnt2[cc * m]);
        }
    for (cc = 1; cc ≤ Next; cc++) {
        u64 val = cc ≤ Wb ? cnt[cc * m] : cnt2[cc * m];
        if (val) printf("open_%dx%d=%llu%s\n", m, cc, val, cc > Wb ? "_ (SpMV)" : "");
    }
    fprintf(stderr, "%s\n", good ? "SpMV_matches_direct" : "SpMV_MISMATCH");
}

```

29. Block SpMV: compute several primes' residues in ONE pass over the edge table. The per-prime state vectors are interleaved ($v[i * B + b]$), so a cache line serves several primes at once; the edges (sorted by dst) are summed per- dst in a u64 scratch and reduced once, not once per edge. This amortizes the dominant edge bandwidth across the whole prime batch – the big win when the table is RAM-resident.

⟨Subroutines 4⟩ +=

```

void spmv_run_batch(int Nto, u64 *pr, int B) { int cc, b;
    long i;
    if (c0_g < 0 ∨ Plevs ≤ 0) {
        for (cc = 1; cc ≤ Nto; cc++)
            if (cc ≤ direct_valid_col_g) {
                printf("%d", cc);
                for (b = 0; b < B; b++) printf("%11u", (unsigned long long)(u64)(cnt[cc * m] % pr[b]));
                printf("\n");
            }
        return;
    }
    u32 *cnt2b = xalloc((size_t)(1 ≪ 16) * B, sizeof(u32));
    u32 *v = xmalloc((size_t) nstate[0] * B * sizeof(u32));
    for (i = 0; i < nstate[0]; i++) {
        u128 sd = seedv[i];
        for (b = 0; b < B; b++) v[i * B + b] = (u32)(sd % pr[b]);
    }
    int basecol = c0_g + 1;
    static Edge *lbuf = 0;
    static long lbufcap = 0; while (basecol ≤ Nto) { int L; for (L = 0; L < Plevs; L++) { int
        abscol = basecol + L/m, substep = L % m;
        long e;
        for (e = 0; e < ncomp[L]; e++) {
            int idx = abscol * m + substep + 1 + comps[L][e].delta;
            long sc = comps[L][e].src;
            u64 mult = comps[L][e].mult;
            if (idx ≥ 0 ∧ idx < (1 ≪ 16))
                for (b = 0; b < B; b++) {
                    size_t o = (size_t) idx * B + b;
                    cnt2b[o] = (u32)(((u64) cnt2b[o] + (u64) v[(size_t) sc * B + b] * mult) % pr[b]);
                }
        }
        u32 *vn = xalloc((size_t) nstate[L + 1] * B, sizeof(u32));
        long ne = nedge[L]; /* read the level's edges (served from page cache if resident), then apply in
            parallel: split at dst boundaries so threads write disjoint vn */
        if (lbufcap < ne) {
            lbufcap = ne;
            lbuf = xrealloc(lbuf, lbufcap * sizeof(Edge));
        }
        fseeko(tbl_fp, lev_edge_off[L], SEEK_SET);
        {
            long rd = 0;
            while (rd < ne) {
                long ch = ne - rd;

```

```

    if (ch > (1_L << 24)) ch = 1_L << 24;
    if ((long) fread(lbuf + rd, sizeof(Edge), ch, tbl_fp) != ch) {
        fprintf(stderr, "table_read_error\n");
        exit(3);
    }
    rd += ch;
}
}
{ int NT = omp_get_max_threads(), tt;
static long pbnd[NTMAX + 1];
pbnd[0] = 0;
pbnd[NT] = ne;
for (tt = 1; tt < NT; tt++) {
    long g = (long)((double) tt * ne / NT);
    while (g > 0 & g < ne & lbuf[g].dst == lbuf[g - 1].dst) g++;
    pbnd[tt] = g;
}
#pragma omp parallel for schedule (dynamic, 1)
for (tt = 0; tt < NT; tt++) {
    long e2, cur = -1;
    u64 acc[64];
    int bb;
    for (e2 = pbnd[tt]; e2 < pbnd[tt + 1]; e2++) {
        long dst = lbuf[e2].dst, src = lbuf[e2].src;
        u64 c = lbuf[e2].c;
        if (dst != cur) {
            if (cur >= 0)
                for (bb = 0; bb < B; bb++) vn[(size_t) cur * B + bb] = (u32)(acc[bb] % pr[bb]);
            cur = dst;
            for (bb = 0; bb < B; bb++) acc[bb] = 0;
        }
        for (bb = 0; bb < B; bb++) acc[bb] += (u64) v[(size_t) src * B + bb] * c;
    }
    if (cur >= 0)
        for (bb = 0; bb < B; bb++) vn[(size_t) cur * B + bb] = (u32)(acc[bb] % pr[bb]);
}
} free(v);
v = vn; } basecol += period_g; } free(v);
for (cc = 1; cc <= Nto; cc++) {
    printf("%d", cc);
    for (b = 0; b < B; b++) {
        u64 r = cc <= direct_valid_col_g ? (u64)(cnt[cc * m] % pr[b]) : (u64) cnt2b[(size_t)(cc * m) * B + b];
        printf("%11u", (unsigned long long) r);
    }
    printf("\n");
}
free(cnt2b); }

```

30. Compressed edge tables (scheme A: CSC-style delta + varint). Each level's coalesced, *dst*-sorted edges are encoded as a byte stream of groups $\langle \text{dst-delta}, \text{gsize}, (\text{src-delta}, \text{c})^{\text{gsize}} \rangle$ (all LEB128 varints); per-group boundaries let us store a handful of split points (byte offset + running dst) so the stream decodes in parallel. Typically ~ 3 bytes/edge vs 16, so a 100 GB table shrinks to ~ 25 GB – it then fits in RAM, turning the disk-bound SpMV into a cache-bound one, and cuts memory traffic $\sim 5\times$.

$\langle \text{Globals } 3 \rangle + \equiv$

```

unsigned char *cedge[MAXLEV];
long ced_len[MAXLEV];
long *csuff[MAXLEV];
long *csdst[MAXLEV];
int csn[MAXLEV];

```

31. $\langle \text{Subroutines 4} \rangle + \equiv$

```

static int vput(unsigned char *p, u64 x)
{
    int n = 0;
    while (x ≥ #80) {
        p[n++] = (x & #7f) | #80;
        x >>= 7;
    }
    p[n++] = (unsigned char) x;
    return n;
}

static u64 vget(unsigned char **pp)
{
    unsigned char *p = *pp;
    u64 x = 0;
    int sh = 0;
    unsigned char b;
    do {
        b = *p++;
        x |= (u64)(b & #7f) << sh;
        sh += 7;
    } while (b & #80);
    *pp = p;
    return x;
}

void compress_table(const char *in, const char *out)
{
    FILE *f = fopen(in, "rb");
    if (!f) {
        fprintf(stderr, "cannot open %s\n", in);
        exit(1);
    }
    int hdr[6];
    if (fread(hdr, sizeof(int), 6, f) ≠ 6) {
        fprintf(stderr, "bad table\n");
        exit(1);
    }
    m = hdr[0];
    period_g = hdr[1];
    c0_g = hdr[2];
    Plevs = hdr[3];
    recend_col_g = hdr[4];
    direct_valid_col_g = hdr[5];
    fread(&seedn, sizeof(long), 1, f);
    seedv = xmalloc(seedn * sizeof (u128));
    fread(seedv, sizeof (u128), seedn, f);
    fread(nstate, sizeof(long), Plevs + 1, f);
    FILE *g = fopen(out, "wb");
    fwrite(hdr, sizeof(int), 6, g);
    fwrite(&seedn, sizeof(long), 1, g);
}

```

```

fwrite(seedv, sizeof (u128), seedn, g);
fwrite(nstate, sizeof (long), Plevs + 1, g);
long SPLIT = 1_L << 21;
Edge *eb = 0;
long ebcap = 0;
unsigned char *ob = 0;
long obcap = 0;
int L;
for (L = 0; L < Plevs; L++) {
    long ne;
    fread(&ne, sizeof (long), 1, f);
    if (ebcap < ne) {
        ebcap = ne;
        eb = xrealloc(eb, ebcap * sizeof (Edge));
    }
    if (ne) fread(eb, sizeof (Edge), ne, f);
    if (obcap < ne * 6 + 1024) {
        obcap = ne * 6 + 1024;
        ob = xrealloc(ob, obcap);
    }
    long olen = 0, prev_dst = -1, since = 0, e2 = 0;
    int scap = (int)(ne/SPLIT) + 2;
    long *soff = xmalloc(scap * sizeof (long));
    long *sdst = xmalloc(scap * sizeof (long));
    int sn = 0;
    while (e2 < ne) {
        if (e2 == 0 ∨ since ≥ SPLIT) {
            soff[sn] = olen;
            sdst[sn] = prev_dst;
            sn++;
            since = 0;
        }
        long dst = eb[e2].dst;
        olen += vput(ob + olen, (u64)(dst - prev_dst));
        long j = e2;
        while (j < ne ∧ eb[j].dst == dst) j++;
        long gsize = j - e2;
        olen += vput(ob + olen, (u64) gsize);
        long prev_src = 0, k;
        for (k = e2; k < j; k++) {
            olen += vput(ob + olen, (u64)(eb[k].src - prev_src));
            prev_src = eb[k].src;
            olen += vput(ob + olen, eb[k].c);
        }
        prev_dst = dst;
        since += gsize;
        e2 = j;
    }
    fwrite(&olen, sizeof (long), 1, g);

```

```

    fwrite(&sn, sizeof(int), 1, g);
    fwrite(soff, sizeof(long), sn, g);
    fwrite(sdst, sizeof(long), sn, g);
    if (olen) fwrite(ob, 1, olen, g);
    free(soff);
    free(sdst);
    long ncp;
    fread(&ncp, sizeof(long), 1, f);
    Comp *cp = xmalloc((ncp + 1) * sizeof(Comp));
    if (ncp) fread(cp, sizeof(Comp), ncp, f);
    fwrite(&ncp, sizeof(long), 1, g);
    if (ncp) fwrite(cp, sizeof(Comp), ncp, g);
    free(cp);
}

int nc;
fread(&nc, sizeof(int), 1, f);
memset(cnt, 0, sizeof(cnt));
fread(cnt, sizeof(u128), nc, f);
fwrite(&nc, sizeof(int), 1, g);
fwrite(cnt, sizeof(u128), nc, g);
free(eb);
free(ob);
fclose(f);
fclose(g);
fprintf(stderr, "compressed %s -> %s\n", in, out);
}

int load_ctables(const char *path)
{
    FILE *f = fopen(path, "rb");
    if (!f) return 0;
    int L, hdr[6];
    if (fread(hdr, sizeof(int), 6, f) != 6) return 0;
    m = hdr[0];
    period_g = hdr[1];
    c0_g = hdr[2];
    Plevs = hdr[3];
    recend_col_g = hdr[4];
    direct_valid_col_g = hdr[5];
    fread(&seedn, sizeof(long), 1, f);
    seedv = xmalloc(seedn * sizeof(u128));
    fread(seedv, sizeof(u128), seedn, f);
    fread(nstate, sizeof(long), Plevs + 1, f);
    for (L = 0; L < Plevs; L++) {
        fread(&ced_len[L], sizeof(long), 1, f);
        fread(&csn[L], sizeof(int), 1, f);
        csloff[L] = xmalloc((csn[L] + 1) * sizeof(long));
        csdst[L] = xmalloc((csn[L] + 1) * sizeof(long));
        fread(csloff[L], sizeof(long), csn[L], f);
        fread(csdst[L], sizeof(long), csn[L], f);
        cedge[L] = xmalloc(ced_len[L] + 16);
    }
}

```

```

    fread(cedge[L], 1, ced_len[L], f);
    fread(&ncomp[L], sizeof(long), 1, f);
    comps[L] = xmalloc((ncomp[L] + 1) * sizeof(Comp));
    fread(comps[L], sizeof(Comp), ncomp[L], f);
}
int nc;
fread(&nc, sizeof(int), 1, f);
memset(cnt, 0, sizeof(cnt));
fread(cnt, sizeof(u128), nc, f);
fclose(f);
return 1;
}
void spmv_run_batch_c(int Nto, u64 *pr, int B) { int cc, b;
    long i;
    if (c0_g < 0 ∨ Plevs ≤ 0) {
        for (cc = 1; cc ≤ Nto; cc++)
            if (cc ≤ direct_valid_col_g) {
                printf("%d", cc);
                for (b = 0; b < B; b++) printf("_%11u", (unsigned long long)(u64)(cnt[cc * m] % pr[b]));
                printf("\n");
            }
        return;
    }
    u32 *cnt2b = xalloc((size_t)(1 ≪ 16) * B, sizeof(u32));
    u32 *v = xmalloc((size_t)nstate[0] * B * sizeof(u32));
    for (i = 0; i < nstate[0]; i++) {
        u128 sd = seedv[i];
        for (b = 0; b < B; b++) v[i * B + b] = (u32)(sd % pr[b]);
    }
    int basecol = c0_g + 1; while (basecol ≤ Nto) { int L; for (L = 0; L < Plevs; L++) { int
        abscol = basecol + L/m, substep = L % m;
    long e;
    for (e = 0; e < ncomp[L]; e++) {
        int idx = abscol * m + substep + 1 + comps[L][e].delta;
        long sc = comps[L][e].src;
        u64 mult = comps[L][e].mult;
        if (idx ≥ 0 ∧ idx < (1 ≪ 16))
            for (b = 0; b < B; b++) {
                size_t o = (size_t)idx * B + b;
                cnt2b[o] = (u32)((u64)cnt2b[o] + (u64)v[(size_t)sc * B + b] * mult) % pr[b];
            }
    }
    u32 *vn = xalloc((size_t)nstate[L + 1] * B, sizeof(u32));
    int NT = omp_get_max_threads(), tt, SN = csn[L];
#pragma omp parallel for schedule(dynamic, 1)
    for (tt = 0; tt < NT; tt++) {
        int s0 = (int)((long)tt * SN/NT), s1 = (int)((long)(tt + 1) * SN/NT);
        if (s1 > SN) s1 = SN;
        if (s0 < s1) {

```

```

unsigned char *p = cedge[L] + csoff[L][s0];
unsigned char *endp = cedge[L] + (s1 < SN ? csoff[L][s1] : ced_len[L]);
long prev_dst = csdst[L][s0];
u64 acc[64];
int bb;
while (p < endp) {
    long dst = prev_dst + (long) vget(&p);
    long gsize = (long) vget(&p);
    for (bb = 0; bb < B; bb++) acc[bb] = 0;
    long prev_src = 0, gg;
    for (gg = 0; gg < gsize; gg++) {
        long src = prev_src + (long) vget(&p);
        prev_src = src;
        u64 c = vget(&p);
        for (bb = 0; bb < B; bb++) acc[bb] += (u64) v[(size_t) src * B + bb] * c;
    }
    for (bb = 0; bb < B; bb++) vn[(size_t) dst * B + bb] = (u32)(acc[bb] % pr[bb]);
    prev_dst = dst;
}
}
}
free(v);
v = vn; } basecol += period_g; } free(v);
for (cc = 1; cc ≤ Nto; cc++) {
    printf("%d", cc);
    for (b = 0; b < B; b++) {
        u64 r = cc ≤ direct_valid_col_g ? (u64)(cnt[cc*m] % pr[b]) : (u64) cnt2b[(size_t)(cc*m)*B+b];
        printf("%11u", (unsigned long long) r);
    }
    printf("\n");
}
free(cnt2b); }

```

32. Main. No arguments: self-checks (direct sweep vs known values, and SpMV vs direct). *dualhamm WbN*: build to column W_b , then extend to column N by SpMV.

$\langle \text{Main } 32 \rangle \equiv$

```

void spmv_run(int Nto,int crosscheck);
void spmv_run_batch(int Nto,u64 *pr,int B);
void compress_table(const char *,const char *);
int load_ctypeables(const char *);
void spmv_run_batch_c(int Nto,u64 *pr,int B);
int build_extract(int mm,int Wb);
int main(int argc,char *argv[])
{
    start_mem_watcher();
    if (argc  $\geq$  5  $\wedge$   $\neg$ strcmp(argv[1], "build")) {
        /* build m Wb file [ckpt] : build+extract, dump tables */
        stop_after_record = 1;
        stream_edges = 1;
        snprintf(rec_path, sizeof rec_path, "%s", argv[4]);
        if (argc  $\geq$  6) ckpt_path = argv[5];
        if (argc  $\geq$  7) ckpt_every = atoi(argv[6]);
        build_extract(atoi(argv[2]), atoi(argv[3]));
        dump_tables(argv[4]);
        return 0;
    }
    if (argc  $\equiv$  5  $\wedge$   $\neg$ strcmp(argv[1], "resume")) { /* resume ckpt Wb file : continue a build */
        stop_after_record = 1;
        stream_edges = 1;
        snprintf(rec_path, sizeof rec_path, "%s", argv[4]);
        resume_from = load_ckpt(argv[2]);
        ckpt_path = argv[2];
        build_extract(m, atoi(argv[3]));
        dump_tables(argv[4]);
        return 0;
    }
    if (argc  $\geq$  4  $\wedge$   $\neg$ strcmp(argv[1], "run")) {
        /* run file Nto [prime] : load tables, SpMV (mod prime if given) */
        if ( $\neg$ load_tables(argv[2])) {
            fprintf(stderr, "cannot_load_%s\n", argv[2]);
            return 1;
        }
        if (argc  $\geq$  5) MODP = strtoull(argv[4], 0, 10);
        spmv_run(atoi(argv[3]), 0);
        return 0;
    }
    if (argc  $\geq$  5  $\wedge$   $\neg$ strcmp(argv[1], "runb")) {
        /* runb file Nto p1 p2 ... : one table pass, many primes */
        if ( $\neg$ load_tables(argv[2])) {
            fprintf(stderr, "cannot_load_%s\n", argv[2]);
            return 1;
        }
    }
    int B = argc - 4;
    if (B > 64) B = 64;

```

```

    static u64 pr[64];
    int b;
    for (b = 0; b < B; b++) pr[b] = strtoull(argv[4 + b], 0, 10);
    spmv_run_batch(atoi(argv[3]), pr, B);
    return 0;
}
if (argc == 4 ∧ ¬strcmp(argv[1], "compress")) {
    compress_table(argv[2], argv[3]);
    return 0;
}
if (argc ≥ 5 ∧ ¬strcmp(argv[1], "runbc")) {
    /* runbc cfile Nto p1 p2 ... : compressed table, batched */
    if (¬load_ctype(argv[2])) {
        fprintf(stderr, "cannot load %s\n", argv[2]);
        return 1;
    }
    int B = argc - 4;
    if (B > 64) B = 64;
    static u64 pr[64];
    int b;
    for (b = 0; b < B; b++) pr[b] = strtoull(argv[4 + b], 0, 10);
    spmv_run_batch_c(atoi(argv[3]), pr, B);
    return 0;
}
if (argc == 4) {
    run_periodic(atoi(argv[1]), atoi(argv[2]), atoi(argv[3]));
    return 0;
}
/* A light smoke test: m=4 is tiny, and by transposition it also pins the m=6 and m=7 values
   (4xc=cx4) without their million-state sweeps; one m=5 strip, built wide enough to record,
   exercises both the direct and the SpMV paths. */
struct {
    int m, Wb, Next, c;
    u64 e;
} chk[] = {
    {4, 9, 9, 6, 744ULL}, /* 4x6 = 6x4 */
    {4, 9, 9, 7, 6378ULL}, /* 4x7 = 7x4 */
    {5, 18, 22, 8, 18061054ULL}, /* m=5, direct column */
    {5, 18, 22, 12, 3611823644006ULL}, /* m=5, direct seam column (recend+1) */
    {5, 18, 22, 15, 24535910156176100ULL}, /* m=5, SpMV extrapolation (past recend+1) */
    {0, 0, 0, 0, 0};
int i, bad = 0, lm = 0, lw = 0, ln = 0;
for (i = 0; chk[i].m; i++) {
    if (chk[i].m ≠ lm ∨ chk[i].Wb ≠ lw ∨ chk[i].Next ≠ ln) {
        run_periodic(chk[i].m, chk[i].Wb, chk[i].Next);
        lm = chk[i].m;
        lw = chk[i].Wb;
        ln = chk[i].Next;
    }
    int c = chk[i].c;
    u64 g = (c ≤ direct_valid_col_g ? cnt[c * m] : cnt2[c * m]), e = red(chk[i].e);
    printf("open %dx%d = %llu exp %llu %s\n", chk[i].m, c, g, e, g == e ? "OK" : "FAIL");
}

```

```

    if (g ≠ e) bad++;
}
printf("%s\n", bad ? "SOME_FAILED" : "ALL_OK");
return bad ? 1 : 0;
}

```

This code is used in section 1.

__int128: [2](#).
_exit: [5](#), [6](#).
_FILE_OFFSET_BITS: [1](#).
_GNU_SOURCE: [1](#).
_SC_PAGE_SIZE: [6](#).
_SC_PHYS_PAGES: [6](#).
A: [11](#), [12](#), [13](#).
a: [4](#), [11](#), [12](#), [13](#), [16](#).
abscol: [27](#), [29](#), [31](#).
acc: [29](#), [31](#).
add_derived: [8](#), [9](#), [16](#).
apexnew: [13](#), [15](#), [16](#).
apexpos: [9](#).
arg: [6](#), [12](#).
argc: [32](#).
argv: [32](#).
atoi: [32](#).
atol: [6](#).
avail: [6](#).
B: [11](#), [12](#), [13](#), [29](#), [31](#), [32](#).
b: [4](#), [11](#), [12](#), [13](#), [16](#), [29](#), [31](#), [32](#).
bad: [32](#).
base: [6](#), [12](#).
basecol: [27](#), [29](#), [31](#).
bb: [29](#), [31](#).
bmate: [13](#), [14](#), [16](#).
bnd: [12](#).
buf: [27](#).
bufcap: [27](#).
build_bmate: [14](#), [16](#).
build_board: [7](#), [23](#).
build_extract: [26](#), [32](#).
c: [7](#), [12](#), [13](#), [23](#), [29](#), [31](#), [32](#).
calloc: [4](#).
cand_c: [23](#), [24](#).
cand_p: [23](#), [24](#).
cb: [11](#).
cc: [7](#), [12](#), [16](#), [26](#), [28](#), [29](#), [31](#).
ced_len: [30](#), [31](#).
cedge: [30](#), [31](#).
cg: [6](#).
cgroup_mem_limit: [6](#).
ch: [29](#).
CH: [27](#).
chk: [32](#).
chunk: [27](#).

ckpt_every: [13](#), [23](#), [32](#).
ckpt_path: [13](#), [19](#), [32](#).
cnt: [13](#), [16](#), [17](#), [19](#), [21](#), [23](#), [26](#), [28](#), [29](#), [31](#), [32](#).
cnt2: [23](#), [26](#), [27](#), [28](#), [32](#).
cnt2b: [29](#), [31](#).
colfp: [24](#).
colfp_g: [13](#), [19](#), [23](#), [24](#).
Comp: [12](#), [13](#), [17](#), [21](#), [31](#).
completion_mp: [8](#), [9](#), [16](#).
compress_table: [31](#), [32](#).
comps: [12](#), [13](#), [21](#), [27](#), [29](#), [31](#).
cp: [31](#).
critical: [17](#).
crosscheck: [26](#), [32](#).
csdst: [30](#), [31](#).
csn: [30](#), [31](#).
csoff: [30](#), [31](#).
cur: [29](#).
curkl: [12](#), [13](#), [16](#), [19](#), [23](#).
curkp: [12](#), [13](#), [16](#), [19](#), [23](#).
curlen: [12](#).
curoff: [12](#), [13](#), [16](#), [19](#), [23](#).
curpos: [12](#).
curw: [12](#), [13](#), [16](#), [19](#), [23](#).
cycle: [8](#), [9](#), [16](#).
c0: [23](#), [24](#), [25](#), [28](#).
c0_g: [13](#), [21](#), [23](#), [24](#), [26](#), [27](#), [29](#), [31](#).
d: [11](#), [12](#), [27](#).
deg: [16](#).
delta: [13](#), [17](#), [27](#), [29](#), [31](#).
direct_cov_g: [13](#), [23](#), [26](#).
direct_valid_col_g: [13](#), [21](#), [23](#), [25](#), [26](#), [29](#), [31](#), [32](#).
done: [27](#).
dst: [12](#), [13](#), [14](#), [27](#), [29](#), [30](#), [31](#).
dualham: [32](#).
dump_tables: [21](#), [32](#).
dynamic: [12](#), [16](#), [27](#), [29](#), [31](#).
e: [6](#), [27](#), [29](#), [31](#), [32](#).
eb: [12](#), [31](#).
ebase: [12](#).
ebcap: [31](#).
ecmp: [12](#), [13](#).
Edge: [12](#), [13](#), [21](#), [27](#), [29](#), [31](#).
edges: [12](#), [13](#), [21](#), [27](#).
emit: [12](#), [16](#).
enb: [12](#).

endp: [31](#).
exit: [4](#), [12](#), [21](#), [23](#), [25](#), [27](#), [29](#), [31](#).
e2: [27](#), [29](#), [31](#).
f: [6](#), [19](#), [21](#), [31](#).
fclose: [6](#), [19](#), [21](#), [31](#).
ferror: [12](#), [21](#), [23](#).
fflush: [6](#), [21](#).
fgets: [6](#).
fopen: [6](#), [19](#), [21](#), [23](#), [31](#).
fprintf: [4](#), [6](#), [12](#), [19](#), [21](#), [23](#), [24](#), [25](#), [26](#), [27](#),
[28](#), [29](#), [31](#), [32](#).
fr: [3](#), [7](#), [15](#).
fread: [19](#), [21](#), [27](#), [29](#), [31](#).
free: [27](#), [29](#), [31](#).
frn: [9](#).
frnew: [15](#), [16](#).
froid: [15](#).
frontier_before: [2](#), [7](#), [15](#), [23](#).
fscanf: [6](#).
fseek: [21](#).
fseeko: [21](#), [27](#), [29](#).
ftell: [23](#).
ftello: [21](#).
fwrite: [12](#), [19](#), [21](#), [23](#), [31](#).
g: [6](#), [27](#), [29](#), [31](#), [32](#).
getenv: [6](#), [23](#).
gg: [31](#).
good: [26](#), [28](#).
gsize: [31](#).
h: [23](#).
hdr: [21](#), [23](#), [31](#).
hi: [12](#).
hn: [12](#).
hp: [12](#).
hpos: [12](#).
i: [9](#), [12](#), [14](#), [15](#), [16](#), [19](#), [23](#), [26](#), [29](#), [31](#), [32](#).
idx: [12](#), [27](#), [29](#), [31](#).
ifrb: [3](#), [7](#), [15](#).
ifrnew: [15](#).
in: [31](#).
in_ncur: [15](#).
in_ncur_g: [12](#), [15](#), [18](#).
inF: [7](#).
i2: [27](#).
j: [9](#), [12](#), [31](#).
k: [7](#), [9](#), [12](#), [31](#).
kb: [19](#).
kbase: [12](#).
KC: [3](#), [7](#).
kcap: [11](#), [12](#).
key: [9](#), [12](#), [23](#).
keycmp: [12](#).

keyof: [9](#), [16](#), [23](#).
kk: [23](#).
kl: [23](#).
kp: [11](#), [12](#).
KR: [3](#), [7](#).
kuse: [11](#), [12](#), [19](#).
L: [21](#), [23](#), [27](#), [29](#), [31](#).
l: [11](#), [12](#).
la: [12](#).
lb: [12](#).
lbuf: [29](#).
lbufcap: [29](#).
len: [10](#), [11](#), [12](#).
lev_edge_off: [13](#), [21](#), [27](#), [29](#).
lim: [6](#).
lm: [32](#).
ln: [32](#).
lo: [12](#).
load_ckpt: [19](#), [32](#).
load_ctype: [31](#), [32](#).
load_tables: [21](#), [32](#).
lw: [32](#).
m: [3](#), [32](#).
main: [32](#).
malloc: [4](#).
mate: [8](#), [9](#), [16](#), [23](#).
max: [6](#).
MAXF: [3](#), [8](#), [13](#), [15](#), [16](#), [23](#).
MAXLEV: [3](#), [13](#), [23](#), [30](#).
mem_cap_bytes: [5](#), [6](#).
mem_watcher: [6](#).
memcmp: [11](#), [12](#).
memcpy: [12](#), [23](#).
memory: [6](#).
memset: [21](#), [23](#), [26](#), [31](#).
merge_pools: [12](#).
mid: [12](#).
mm: [6](#), [19](#), [23](#), [26](#), [32](#).
mmap: [5](#).
MODP: [13](#), [27](#), [32](#).
mp: [9](#), [16](#), [17](#).
mt: [12](#).
mult: [13](#), [17](#), [27](#), [29](#), [31](#).
n: [3](#), [4](#), [31](#).
NB: [3](#), [7](#), [15](#).
nbr: [15](#), [16](#).
nc: [21](#), [31](#).
ncomp: [12](#), [13](#), [21](#), [27](#), [29](#), [31](#).
nep: [31](#).
ncur: [12](#), [13](#), [15](#), [16](#), [19](#), [23](#).
ND: [3](#), [7](#), [15](#).
ne: [27](#), [29](#), [31](#).

nedge: [12](#), [13](#), [21](#), [27](#), [29](#).
need: [16](#).
Next: [23](#), [28](#), [32](#).
nexts: [19](#).
nk: [16](#).
nl: [6](#), [16](#).
np: [12](#).
nr: [11](#), [12](#).
NRNG: [12](#).
nstate: [12](#), [13](#), [21](#), [23](#), [25](#), [27](#), [29](#), [31](#).
nstate_off: [13](#), [21](#), [23](#).
nt: [12](#).
NT: [12](#), [27](#), [29](#), [31](#).
NTMAX: [11](#), [12](#), [27](#), [29](#).
Nto: [26](#), [27](#), [29](#), [31](#), [32](#).
o: [12](#), [29](#), [31](#).
ob: [31](#).
obcap: [31](#).
off: [10](#), [11](#), [12](#).
off_t: [21](#).
ok: [16](#).
okl: [16](#).
olen: [31](#).
omate: [14](#), [16](#).
omp: [8](#), [12](#), [13](#), [16](#), [17](#), [27](#), [29](#), [31](#).
omp_get_max_threads: [12](#), [27](#), [29](#), [31](#).
omp_get_thread_num: [12](#).
op: [14](#).
out: [31](#).
o2: [12](#).
o2n: [13](#), [14](#), [15](#).
P: [12](#).
p: [4](#), [12](#), [31](#).
parallel: [12](#), [16](#), [27](#), [29](#), [31](#).
path: [6](#), [19](#), [21](#), [31](#).
pbnd: [29](#).
pdst: [14](#).
period: [23](#), [24](#).
period_g: [13](#), [21](#), [23](#), [24](#), [27](#), [29](#), [31](#).
Plevs: [13](#), [21](#), [23](#), [24](#), [25](#), [27](#), [29](#), [31](#).
posS: [13](#), [14](#), [15](#), [16](#).
pp: [12](#), [31](#).
pr: [29](#), [31](#), [32](#).
prev_dst: [31](#).
prev_src: [31](#).
printf: [26](#), [28](#), [29](#), [31](#), [32](#).
pthread_create: [6](#).
pthread_detach: [6](#).
pthread_t: [6](#).
q: [4](#), [7](#), [9](#), [23](#).
qnew: [13](#), [14](#), [15](#), [16](#).
qold: [14](#), [15](#), [16](#).

qsort: [5](#), [12](#).
qsort_r: [12](#).
q2: [12](#).
R: [12](#).
r: [6](#), [7](#), [12](#), [29](#), [31](#).
Ra: [12](#).
ram: [6](#).
Rb: [12](#).
rc: [11](#), [12](#).
rcap: [11](#), [12](#).
rcmp: [11](#).
rcmp_r: [12](#).
rcnt: [12](#).
rd: [29](#).
re: [12](#).
realloc: [4](#).
Rec: [10](#), [11](#), [12](#).
rec: [12](#), [13](#), [15](#), [17](#), [18](#).
rec_fp: [12](#), [13](#), [21](#), [23](#).
rec_path: [12](#), [13](#), [23](#), [32](#).
recap: [12](#).
reccomp_buf: [12](#), [13](#), [17](#).
reccomp_cap: [13](#), [17](#).
reccomp_n: [12](#), [13](#), [15](#), [17](#).
recend: [23](#), [24](#), [25](#).
recend_col_g: [13](#), [21](#), [23](#), [24](#), [25](#), [31](#).
reclev: [12](#), [13](#), [23](#).
recording: [13](#), [23](#), [24](#).
recstart: [23](#), [24](#).
red: [12](#), [13](#), [17](#), [26](#), [27](#), [32](#).
res: [6](#).
resume_from: [23](#), [32](#).
rgcap: [12](#).
rk: [12](#).
rkcap: [12](#).
rkey: [12](#).
rkl: [12](#).
rkuse: [12](#).
rlen: [12](#).
RLESS: [12](#).
RLIMIT_AS: [5](#).
rne: [12](#).
rne2: [12](#).
rr: [7](#), [15](#), [16](#).
rss_bytes: [6](#).
run_periodic: [23](#), [26](#), [32](#).
run_step: [15](#), [23](#).
rw: [12](#).
r2: [12](#).
S: [12](#).
s: [6](#), [7](#), [9](#), [12](#), [15](#), [23](#).
same: [12](#).

samp: [12](#).
Samp: [12](#).
sampcap: [12](#).
sampcmp: [12](#).
save_ckpt: [19](#), [23](#).
sc: [12](#), [29](#), [31](#).
scap: [31](#).
schedule: [12](#), [16](#), [27](#), [29](#), [31](#).
sd: [29](#), [31](#).
sdst: [31](#).
seam_break: [23](#), [24](#), [25](#).
seed_bucket: [23](#).
seedn: [13](#), [21](#), [23](#), [31](#).
seedv: [13](#), [21](#), [23](#), [27](#), [29](#), [31](#).
SEEK_CUR: [21](#).
SEEK_SET: [21](#), [27](#), [29](#).
sh: [31](#).
si: [12](#), [15](#), [16](#), [17](#).
since: [31](#).
sm: [12](#).
sn: [31](#).
SN: [31](#).
snprintf: [6](#), [32](#).
soff: [31](#).
sort_reduce: [12](#), [18](#).
spl: [12](#), [27](#).
SPLIT: [31](#).
spmv_run: [23](#), [26](#), [32](#).
spmv_run_batch: [29](#), [32](#).
spmv_run_batch_c: [31](#), [32](#).
src: [10](#), [12](#), [13](#), [17](#), [27](#), [29](#), [31](#).
start_mem_watcher: [6](#), [32](#).
stderr: [4](#), [6](#), [12](#), [19](#), [21](#), [23](#), [24](#), [25](#), [26](#), [27](#),
[28](#), [29](#), [31](#), [32](#).
STEMP: [13](#), [14](#), [16](#).
stop_after_record: [13](#), [23](#), [32](#).
strchr: [6](#).
strcmp: [6](#), [32](#).
stream_edges: [12](#), [13](#), [21](#), [23](#), [32](#).
stream_spmv: [13](#), [21](#), [27](#).
strlen: [6](#).
strncmp: [6](#).
strrchr: [6](#).
strstr: [6](#).
strtoull: [32](#).
substep: [27](#), [29](#), [31](#).
sw: [12](#).
swept: [25](#).
sysconf: [6](#).
sz: [6](#).
s0: [31](#).
s1: [31](#).
t: [12](#), [23](#).
take: [12](#).
tbl_fp: [13](#), [21](#), [27](#), [29](#).
th: [6](#).
threadprivate: [8](#), [13](#), [16](#).
tkcap: [11](#), [12](#).
tkp: [11](#), [12](#).
tkuse: [11](#), [12](#).
tnr: [11](#), [12](#).
tot: [12](#).
trc: [11](#), [12](#).
trcap: [11](#), [12](#).
tt: [12](#), [27](#), [29](#), [31](#).
u: [7](#).
up: [6](#).
usleep: [6](#).
u128: [2](#), [10](#), [12](#), [13](#), [16](#), [19](#), [21](#), [23](#), [29](#), [31](#).
u32: [2](#), [29](#), [31](#).
u64: [2](#), [12](#), [13](#), [19](#), [23](#), [26](#), [27](#), [28](#), [29](#), [31](#), [32](#).
V: [3](#).
v: [6](#), [7](#), [27](#), [29](#), [31](#).
val: [26](#), [28](#).
vget: [31](#).
vn: [27](#), [29](#), [31](#).
vput: [31](#).
w: [15](#).
Wb: [23](#), [26](#), [28](#), [32](#).
x: [6](#), [12](#), [31](#).
xcalloc: [4](#), [27](#), [29](#), [31](#).
xmalloc: [4](#), [5](#), [12](#), [21](#), [23](#), [27](#), [29](#), [31](#).
xrealloc: [4](#), [12](#), [17](#), [19](#), [27](#), [29](#), [31](#).
z: [23](#).

〈Credit completion [17](#)〉 Used in section [16](#).
 〈Detect and confirm the periodic column [24](#)〉 Used in section [23](#).
 〈Expand every state at cell s [16](#)〉 Used in section [15](#).
 〈Finalize direct coverage and verify periodicity [25](#)〉 Used in section [23](#).
 〈Globals [3](#), [5](#), [8](#), [11](#), [13](#), [30](#)〉 Used in section [1](#).
 〈Iterate the SpMV out to column N [27](#)〉 Used in section [26](#).
 〈Main [32](#)〉 Used in section [1](#).
 〈Report and cross-check [28](#)〉
 〈Sort, reduce, and (if recording) capture the edge table [18](#)〉 Used in section [15](#).
 〈Subroutines [4](#), [6](#), [7](#), [9](#), [12](#), [14](#), [15](#), [19](#), [21](#), [23](#), [26](#), [29](#), [31](#)〉 Used in section [1](#).
 〈Types [2](#), [10](#)〉 Used in section [1](#).

DUALHAM

	Section	Page
Introduction	1	1
Board and frontier	2	2
Mate table, derived edges, key, completion	8	7
Bucket aggregation	10	10
The transfer step	13	18
Checkpointing the build	19	23
Dumping and reloading the extracted tables	21	25
Driver: build, extract the period, then SpMV	23	27
Main	32	42